

Foundations of NLP

Lecture-9

Blue Score, BERT (Bidirectional Encoder Representations)

Recap

- Language modeling
- Recurrent Neural Network and Implementation
- Applications of Recurrent Neural Network
- Language modeling using Long Short-term Memory
- Sequence to Sequence learning
- Attention
- Transformers

Why?

- You covered Transformers for Machine Translation Task.

How would you evaluate the quality of translations?

How to evaluate Machine Translation?

BLEU (Bilingual Evaluation Understudy) ✓

BLEU

Article

Talk

Read

Edit

View history

Tools

From Wikipedia, the free encyclopedia

This article is about the evaluation metric for machine translation. For other uses, see [Bleu \(disambiguation\)](#).

BLEU (bilingual evaluation understudy) is an algorithm for [evaluating](#) the quality of text which has been [machine-translated](#) from one [natural language](#) to another. Quality is considered to be the correspondence between a machine's output and that of a human: "the closer a machine translation is to a professional human translation, the better it is" – this is the central idea behind BLEU.^[1] Invented at [IBM](#) in 2001,^[1] BLEU was one of the first [metrics](#) to claim a high [correlation](#) with human judgements of quality,^{[2][3]} and remains one of the most popular automated and inexpensive metrics.

Scores are calculated for individual translated segments—generally sentences—by comparing them with a set of good quality reference translations. Those scores are then averaged over the whole [corpus](#) to reach an estimate of the translation's overall quality. Intelligibility or grammatical correctness are not taken into account.^[4]

BLEU's output is always a number between 0 and 1. This value indicates how similar the candidate text is to the reference texts, with values closer to 1 representing more similar texts. Few human translations will attain a score of 1, since this would indicate that the candidate is identical to one of the reference translations. For this reason, it is not necessary to attain a score of 1. Because there are more opportunities to match, adding additional reference translations will increase the BLEU score.^[5]

Mathematical definition [[edit](#)]

Basic setup [[edit](#)]

A basic, first attempt at defining the BLEU score would take two arguments: a candidate string \hat{y} and a list of reference strings $(y^{(1)}, \dots, y^{(N)})$. The idea is that $BLEU(\hat{y}; y^{(1)}, \dots, y^{(N)})$ should be close to 1 when \hat{y} is similar to $y^{(1)}, \dots, y^{(N)}$, and close to 0 if not.

As an analogy, the BLEU score is like a language teacher trying to score the quality of a student translation \hat{y} by checking how closely it follows the reference answers $y^{(1)}, \dots, y^{(N)}$.

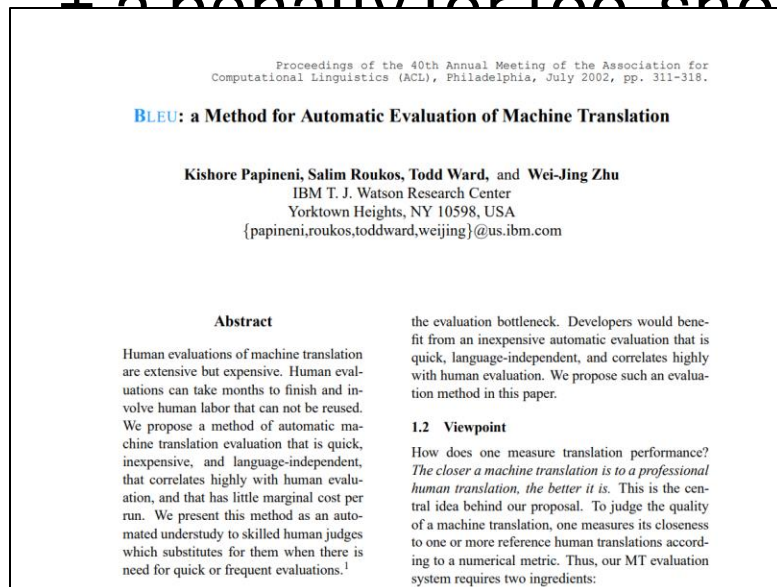
Since in natural language processing, one should evaluate a large set of candidate strings, one must generalize the BLEU score to the case where one has a list of M candidate strings (called a "[corpus](#)") $(\hat{y}^{(1)}, \dots, \hat{y}^{(M)})$, and for each candidate string $\hat{y}^{(i)}$, a list of reference candidate strings $S_i := (y^{(i,1)}, \dots, y^{(i,N_i)})$.

How to evaluate Machine Translation?

BLEU (Bilingual Evaluation Understudy)

BLEU compares the machine-written translation to one or several human-written translation(s), and computes a similarity score based on:

- Geometric mean of n-gram precision (usually for 1, 2, 3 and 4-grams)
- penalty for too short system translations



Source paper: "*BLEU: a Method for Automatic Evaluation of Machine Translation*", Papineni et al, 2002

How to calculate BLEU Score?

French: Le chat est sur le tapis → Source language

Reference: The cat is on the mat → Human-written Translation

Model

MT Output- Candidate: The cat on mat → Machine generated Translation

$$p = \frac{m}{W_t}$$

$$= \frac{4}{4} \quad \text{①}$$

$$m = 4$$
$$W_t = 4$$

$m =$ No of candidate words in the reference

$W_t =$ Total No of words in the candidate

How to calculate BLEU Score? (What can go wrong?)

French: Le chat est sur le tapis

Reference 1: The cat is on the mat

Reference 2: There is a cat on the mat

MT Output- Candidate: the the the the the the the

Reference

$$p = \frac{m}{w_t} = \frac{7}{7} = 1$$

m = 7
w_t = 7
No 00 candidates
words in the
refer

How to calculate BLEU Score?

(NT) The cat is on the mat

the cat the cat on the mat

My Neural Machine Translation model:
generates some output

BLEU Score:



	Count	Count	
the cat	2	1	$P = \frac{7}{7} = 1$
cat the	1	0	
cat on	1	0	
on the	1	1	
the mat	1	1	

Modified =

$\frac{3}{6}$

How to calculate BLEU Score?

Source paper: "*BLEU: a Method for Automatic Evaluation of Machine Translation*", Papineni et al, 2002

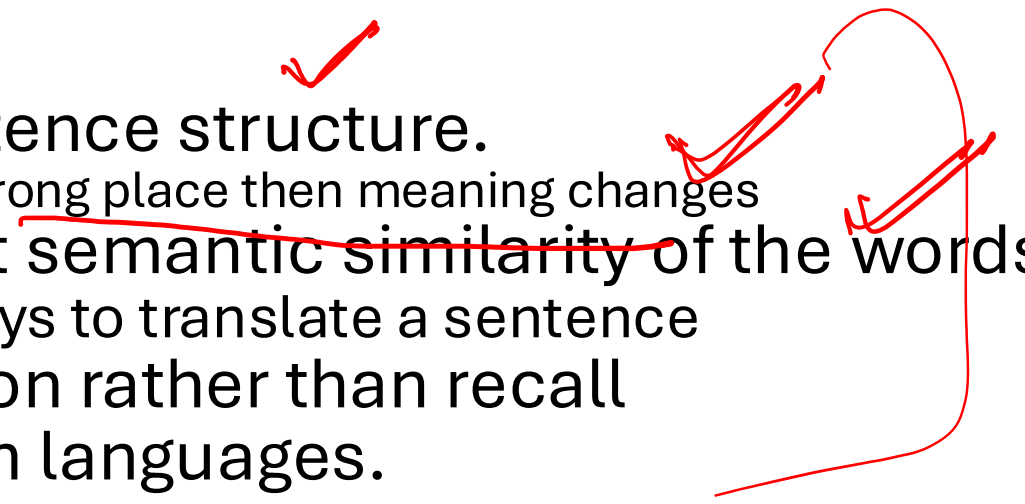
<https://aclanthology.org/P02-1040.pdf>

BLEU Score

Pros:

- Simple and widely used metric
- BLEU is useful but imperfect

Cons:

- Doesn't incorporate sentence structure.
 - If the word is placed in the wrong place then meaning changes
 - Does not take in account semantic similarity of the words
 - There are many valid ways to translate a sentence
 - More focused on precision rather than recall
 - Struggle with non-english languages.
- 

BLEU Score Python Code

```
[ ] ! pip install datasets transformers[sentencepiece]

[ ] from datasets import load_metric

    bleu = load_metric("bleu")
    predictions = [["I", "have", "thirty", "six", "years"]]
    references = [
        ["I", "am", "thirty", "six", "years", "old"], ["I", "am", "thirty", "six"]]
    bleu.compute(predictions=predictions, references=references)

[ ] predictions = [["I", "have", "thirty", "six", "years"]]
    references = [
        ["I", "am", "thirty", "six", "years", "old"], ["I", "am", "thirty", "six"]]
    bleu.compute(predictions=predictions, references=references)

[ ] predictions = [["I", "have", "thirty", "six", "years"]]
    references = [
        ["I", "am", "thirty", "six", "years", "old"], ["I", "am", "thirty", "six"]]
    bleu.compute(predictions=predictions, references=references)

[ ] ! pip install sacrebleu

[ ] sacrebleu = load_metric("sacrebleu")
    # SacreBLEU operates on raw text, not tokens
    predictions = ["I have thirty six years"]
    references = ["I am thirty six years old", "I am thirty six"]
    sacrebleu.compute(predictions=predictions, references=references)
```

https://colab.research.google.com/github/huggingface/notebooks/blob/master/course/videos/bleu_metric.ipynb

Performance of Transformer models- Good results

Table 2: The Transformer achieves better BLEU scores than previous state-of-the-art models on the English-to-German and English-to-French newstest2014 tests at a fraction of the training cost.

Model	BLEU		Training Cost (FLOPs)	
	EN-DE	EN-FR	EN-DE	EN-FR
ByteNet [18]	23.75			
Deep-Att + PosUnk [39]		39.2		$1.0 \cdot 10^{20}$
GNMT + RL [38]	24.6	39.92	$2.3 \cdot 10^{19}$	$1.4 \cdot 10^{20}$
ConvS2S [9]	25.16	40.46	$9.6 \cdot 10^{18}$	$1.5 \cdot 10^{20}$
MoE [32]	26.03	40.56	$2.0 \cdot 10^{19}$	$1.2 \cdot 10^{20}$
Deep-Att + PosUnk Ensemble [39]		40.4		$8.0 \cdot 10^{20}$
GNMT + RL Ensemble [38]	26.30	41.16	$1.8 \cdot 10^{20}$	$1.1 \cdot 10^{21}$
ConvS2S Ensemble [9]	26.36	41.29	$7.7 \cdot 10^{19}$	$1.2 \cdot 10^{21}$
Transformer (base model)	27.3	38.1	$3.3 \cdot 10^{18}$	
Transformer (big)	28.4	41.8	$2.3 \cdot 10^{19}$	

Pretraining

To create domain-specific Language Models (LLMs), you have two main approaches:


- **Training from Scratch:** Building specialized LLMs entirely from domain-specific data.
- **Continual Pre-training:** Enhancing existing LLMs by further pre-training them with domain-specific data.

Pretraining and Fine-tuning

Pretraining can improve NLP applications by serving as parameter initialization.


Step 1: Pretrain (on language modeling)

Lots of text; learn general things!




Step 2: Finetune (on your task)

Not many labels; adapt to the task!




Pretraining for three types of architectures

The neural architecture influences the type of pretraining, and natural use cases.




Encoders

- Gets bidirectional context – can condition on future!
- How do we train them to build strong representations?



**Encoder-
Decoders**

- Good parts of decoders and encoders?
- What's the best way to pretrain them?



Decoders

- Language models! What we've seen so far.
- Nice to generate from; can't condition on future words


Pretraining Encoders

So far, we've looked at language model pretraining. But encoders get **bidirectional context**, so we can't do language modeling!

Idea: replace some fraction of words in the input with a special [MASK] token; predict these words.

$$h_1, \dots, h_T = \text{Encoder}(w_1, \dots, w_T)$$
$$y_i \sim Ah_i + b$$

Only add loss terms from words that are "masked out." If \tilde{x} is the masked version of x , we're learning $p_\theta(x|\tilde{x})$. Called **Masked LM**.



Pre-training






Key ideas in pretraining


- Make sure your model can process large-scale, diverse datasets
- Don't use labeled data (otherwise you can't scale!)
- Compute-aware scaling

Word Structure and sub words

Let's take a look at the assumptions we've made about a language's vocabulary.

We assume a fixed vocab of **tens of thousands of words**, built from the training set. All **novel words** seen at test time are mapped to a single **UNK**.

	word		vocab mapping	embedding
Common words	hat	→	pizza (index)	
	learn	→	tasty (index)	
Variations	taaaaasty	→	UNK (index)	
misspellings	laern	→	UNK (index)	
novel items	Transformerify	→	UNK (index)	

**Brianna** @_parsimonia_ · 24h
Gooooooooo Vibesssssss

Byte pair Encoding Algorithm

Subword modeling in NLP encompasses a wide range of methods for reasoning about structure below the word level. **(Parts of words, characters, bytes.)**

- The dominant modern paradigm is to learn a vocabulary of parts of words (subword tokens).
- At training and testing time, each word is split into a sequence of known subwords.

Byte-pair encoding is a simple, effective strategy for defining a subword vocabulary.

1. **Start with a vocabulary containing only characters and an “end-of-word” symbol.**
2. **Using a corpus of text, find the most common adjacent characters “a,b”; add “ab” as a subword.**
3. **Replace instances of the character pair with the new subword; repeat until desired vocab size.**

Originally used in NLP for machine translation; now similar methods (WordPiece, SentencePiece) are used in pretrained models, like BERT, GPT.

Byte-pair Encoding

corpus	vocabulary
5 l o w _	_, d, e, i, l, n, o, r, s, t, w
2 l o w e s t _	
6 n e w e r _	
3 w i d e r _	
2 n e w _	

The BPE algorithm first counts all pairs of adjacent symbols: the most frequent is the pair `e r` because it occurs in *newer* (frequency of 6) and *wider* (frequency of 3) for a total of 9 occurrences.² We then merge these symbols, treating `er` as one symbol, and count again:

corpus	vocabulary
5 l o w _	_, d, e, i, l, n, o, r, s, t, w, er
2 l o w e s t _	
6 n e w er _	
3 w i d er _	
2 n e w _	






Now the most frequent pair is `er _`, which we merge; our system has learned that there should be a token for word-final `er`, represented as `er_`:


corpus	vocabulary
5 l o w _	_, d, e, i, l, n, o, r, s, t, w, er, er_
2 l o w e s t _	
6 n e w er_	
3 w i d er_	
2 n e w _	

Word Structure and sub words

Common words end up being a part of the subword vocabulary, while rarer words are split into (sometimes intuitive, sometimes not) components.

In the worst case, **words are split into as many subwords as they have characters**

	word		vocab mapping	embedding
Common words	hat	→	hat	
	learn	→	learn	
Variations	taaaaasty	→	taa## aaa## sty	
misspellings	laern	→	la## ern##	
novel items	Transformerify	→	Transformer## ify	



Brianna @_parsimonia_ · 24h
Goooooooood Vibesssssss

Motivation word meaning and Context

Recall the adage we mentioned at the beginning of the course:

“You shall know a word by the company it keeps” (J. R. Firth 1957: 11)

This quote is a summary of **distributional semantics**, and motivated word2vec. But:

“... the complete meaning of a word is always contextual,
and no study of meaning apart from a complete context
can be taken seriously.” (J. R. Firth 1935)

Consider I **record** the **record**: the two instances of **record** mean different things.

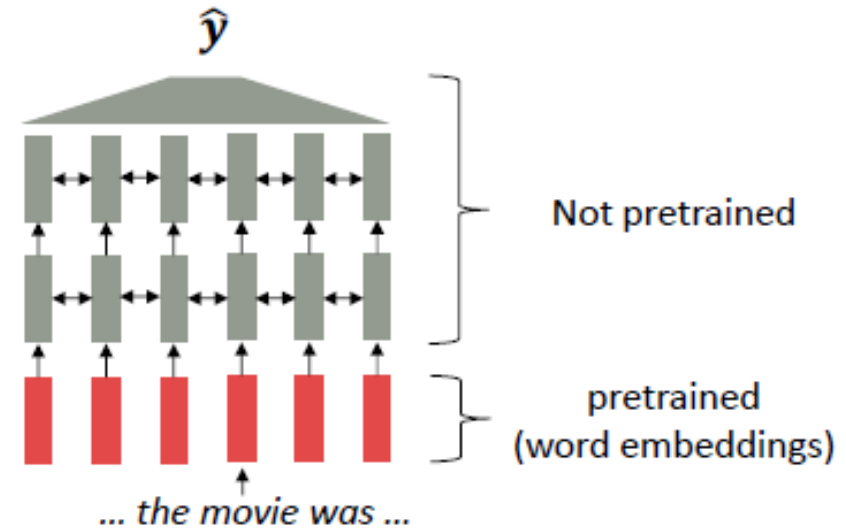
Where were we- Pretrained word embeddings

Circa 2015:

- Start with pretrained word embeddings (no context!)
- Learn how to incorporate context in an LSTM or Transformer while training on the task.

Some issues to think about:

- The training data we have for our **downstream task** (like question answering) must be sufficient to teach all contextual aspects of language.
- Most of the parameters in our network are randomly initialized!



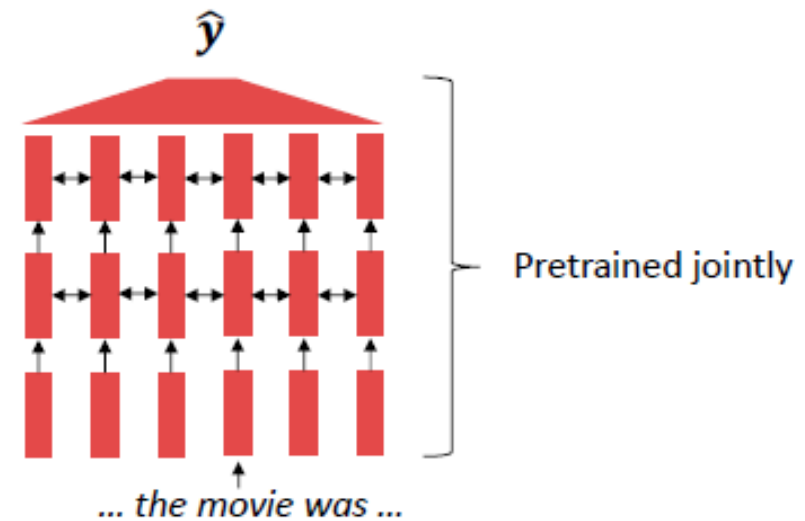
[Recall, *movie* gets the same word embedding, no matter what sentence it shows up in]

How to learn Contextual representations?

Where are we going- Pretraining whole

In modern NLP:

- All (or almost all) parameters in NLP networks are initialized via **pretraining**.
- Pretraining methods hide parts of the input from the model, and train the model to reconstruct those parts.
- This has been exceptionally effective at building strong:
 - **representations of language**
 - **parameter initializations** for strong NLP models.
 - **Probability distributions** over language that we can sample from



[This model has learned how to represent entire sentences through pretraining]

Bidirectional Encoder Representations from Transformers



When predicting words within a sequence, all of the surrounding words can be used to gain contextual information.

Left context

Right context



A man was [MASK] on a river bank.

UNIDIRECTIONAL CONTEXT

When predicting future words, only the previous words can be used to gain contextual information.

Left context

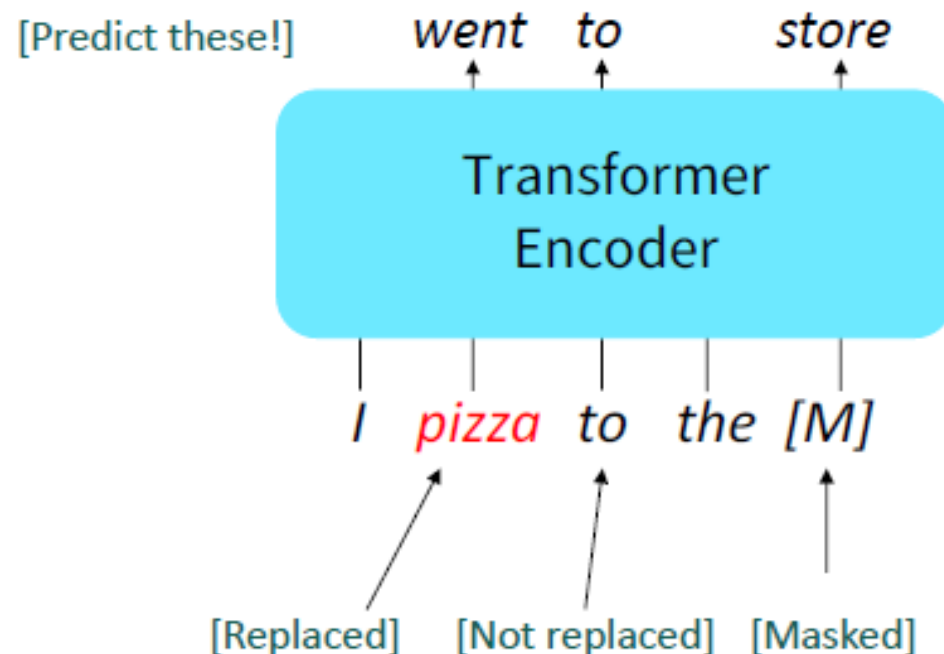


Bidirectional Encoder Representations from Transformers

Devlin et al., 2018 proposed the “Masked LM” objective and released the weights of a pretrained Transformer, a model they labeled BERT.

Some more details about Masked LM for BERT:

- Predict a random 15% of (sub)word tokens.
 - Replace input word with [MASK] 80% of the time
 - Replace input word with a random token 10% of the time
 - Leave input word unchanged 10% of the time (but still predict it!)
- Why? Doesn't let the model get complacent and not build strong representations of non-masked words. (No masks are seen at fine-tuning time!)



Tokenization Algorithm used by BERT

WordPiece is the tokenization algorithm Google developed to pretrain BERT.

WordPiece

Tokenization Process:



<https://huggingface.co/learn/nlp-course/chapter6/6?fw=pt>

BERT Tokenizer

```
[1]: from transformers import BertTokenizer

[2]: tokenizer = BertTokenizer.from_pretrained('bert-base-cased')

[3]: tokenizer.tokenize("This isn't too surprising.")

[3]: ['This', 'isn', "'", 't', 'too', 'surprising', '.']

[4]: tokenizer.tokenize("Encode me!")

[4]: ['En', '##code', 'me', '!']

[5]: tokenizer.tokenize("Snuffleupagus?")

[5]: ['S', '##nu', '##ffle', '##up', '##agu', '##s', '?']

[6]: tokenizer.vocab_size

[6]: 28996
```

Masked Language Modeling

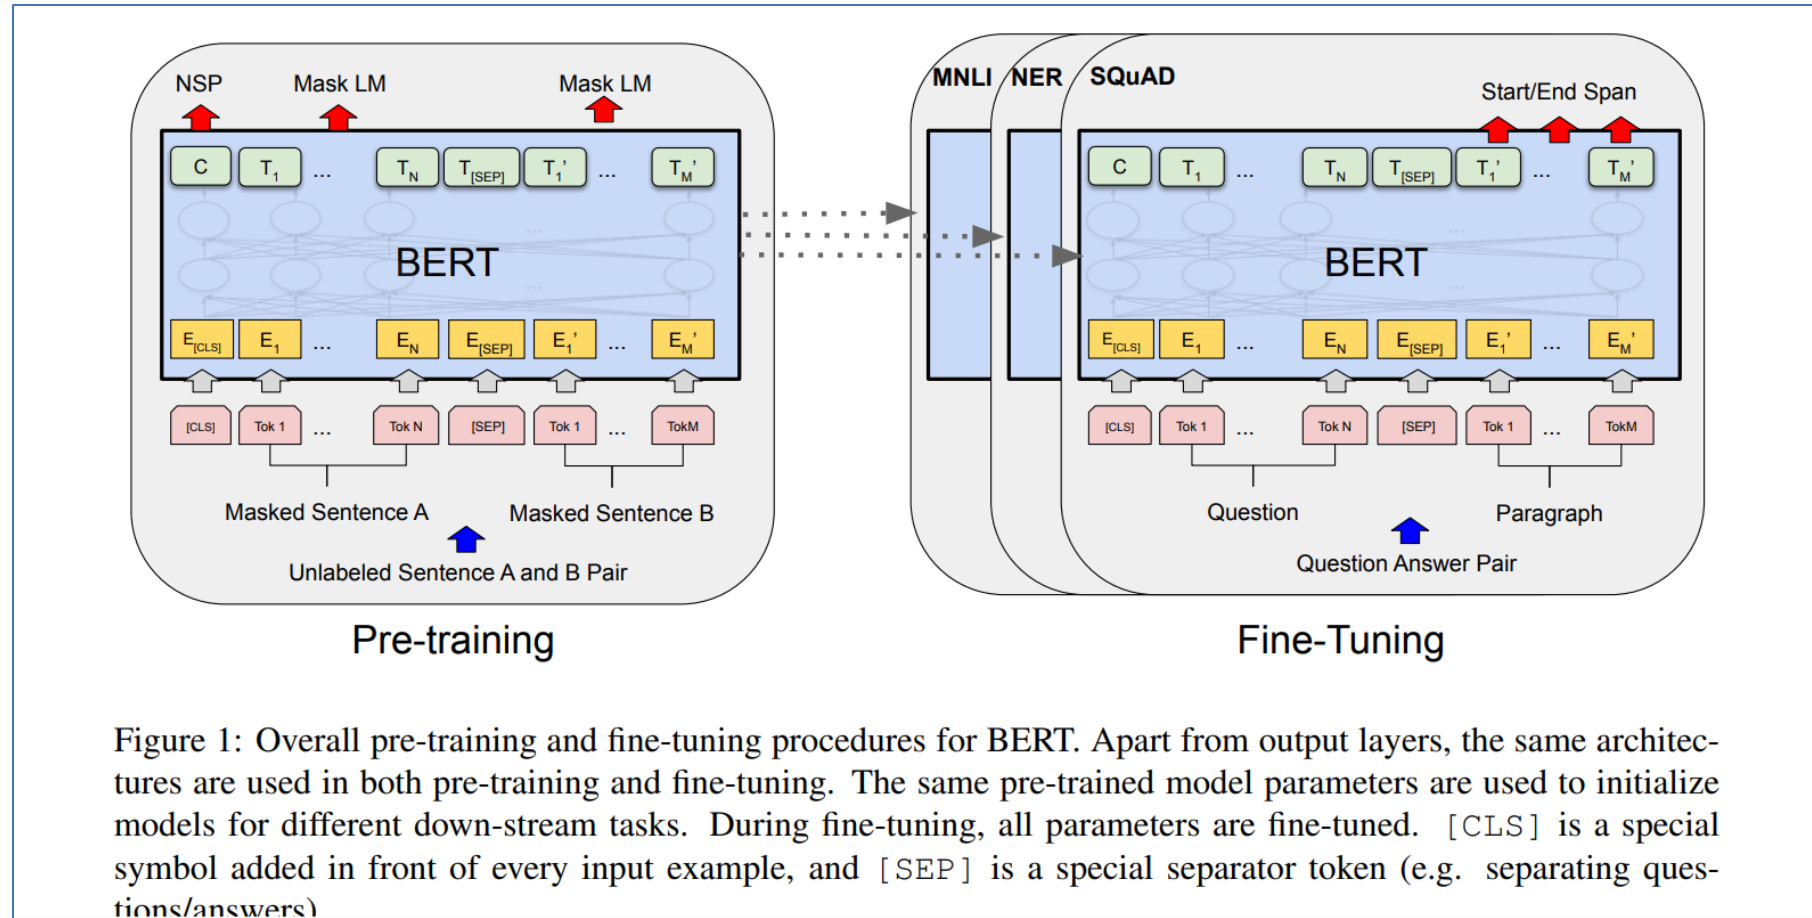
Masked Language Modeling (MLM) (Devlin et al., 2019).

MLM uses unannotated text from a large corpus. Here, the model is presented with a series of sentences from the training corpus where a random sample of tokens from each training sequence is selected for use in the learning task.

Once chosen, a token is used in one of three ways:

- **It is replaced with the unique vocabulary token [MASK].**
- **It is replaced with another token from the vocabulary, randomly sampled based on token unigram probabilities.**
- **It is left unchanged.**

Pretraining and Fine-tuning



<https://arxiv.org/pdf/1810.04805>

BERT Implementation/Code

- [google-research/bert: TensorFlow code and pre-trained models for BERT](#)

Difference from Word2vec

- **Out-of-Vocabulary problem**

- Word2Vec learns embeddings at word level, it can only generate embeddings for words that existed in its training set (aka its “vocabulary space”). This is a major drawback to Word2Vec - that it just doesn’t support Out-of-Vocabulary words.
- BERT learns representations at the subword level, so a BERT model will have a smaller vocabulary space than the number of unique words in its training corpus.
- BERT is able to generate embeddings for words outside of its vocabulary space giving it a near infinite vocabulary.

Difference from Word2vec

- **Embeddings**

- Word2Vec offers pre-trained word embeddings that anyone can use off-the-shelf.
- Word2Vec takes a single word as input and outputs a single vector representation of that word.
- BERT generates contextual embeddings, it takes as input a sequence (usually a sentence) rather than a single word.
- BERT needs to be shown the context that surrounding words provide before it can generate a word embedding.

Difference from Word2vec

- **Probabilistic/Static vs Contextual Embeddings**

- Word2Vec offers pre-trained word embeddings that anyone can use off-the-shelf.
- Word2Vec takes a single word as input and outputs a single vector representation of that word.
- BERT generates contextual embeddings, it takes as input a sequence (usually a sentence) rather than a single word.
- BERT needs to be shown the context that surrounding words provide before it can generate a word embedding.

Acknowledgments

- These slides were adapted from the book SPEECH and LANGUAGE PROCESSING: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition
- Some modifications from CS224N presentations and resources found in the WEB by several scholars.

Reference materials

- <https://vlanc-lab.github.io/mu-nlp-course/teachings/fall2024-AI-schedule.html>
- Lecture notes
- (A) Speech and Language Processing by Daniel Jurafsky and James H. Martin
- (B) Natural Language Processing with Python. (updated edition based on Python 3 and NLTK 3) Steven Bird et al. O'Reilly Media

