

# Foundations of NLP

CS3126

Lecture-5

Probabilistic (N-grams) and Neural Language models

# Recap

- NLP
- Applications
- Regular expressions
- Tokenization
- Stemming
  - Porter Stemmer
- Lemmatization
- Normalization
- Stopwords
- Bag-of-Words
- TF-IDF
- NER
- POS tagging
- Semantics, Distributional semantics, Word2vec

# Minor-I, Do's and Don'ts



- How to prepare for exam?

- **Bad strategy**

- preparing last minute
    - Searching for shortcuts
    - What will come in exam, DMing ("Mam will porter stemmer will be given"? )

- **Good strategy**

- Make use of Office Hours for Doubts
    - Active class participation
    - Use #helpdoubt channel for clarifications.
    - Last minute panic would not help, plan your obstacles well (submission format, links, etc.)!
    - Keep a window for doubts clarifications
    - Have to talk to others, have to check the communication platform (Slack/email)
    - Don't work in isolation, closing eyes, learn about any updates from Slack!

**This will help you maintain reasonable grades throughout (irrespective of difficulty level)**

# Project Groups (First evaluation)

1. Prepare 7-8 slides and upload on the link (to be shared on slack)
  - Motivation
  - Problem Statement
  - Proposed pipeline
  - Timeline
  - Expected Outcome/Application
2. Deadline- 28th September 2024

# Language Model

- **Classic definition-** Probability distribution over sequence of tokens
- Vocabulary  $V$  – a set of tokens!
- A language model  $\mathbf{p}$  assigns each sequence of tokens  $(x_1, \dots, x_L) \in V$ , a probability (a number between 0 and 1),
$$P(x_1, x_2, x_3, \dots, x_L)$$
- The probability intuitively tells us how "good" a sequence of tokens is.

# Language Model

- Consider the probability of four strings in English.
- Here Vocabulary,

$V = \{ate, ball, cheese, mouse, the\}$

- $p(the, mouse, ate, the, cheese) = 0.02 \rightarrow P_1$
- $p(the, cheese, ate, the, mouse) = 0.01 \rightarrow P_2$
- $p(mouse, the, the, cheese, ate) = 0.0001 \rightarrow P_3$

Clearly,  $P_1 > P_2 > P_3$

# Language Model

- LM assigned, "*mouse the the cheese ate*" a very low probability implicitly because it's **ungrammatical (syntactic knowledge)**.
- The LM should assign *the mouse ate the cheese* higher probability than *the cheese ate the mouse* implicitly because of,
  - **world knowledge**: both sentences are the same syntactically,
  - they differ in **semantic plausibility**

# Where do you see language models?

## Google Search system





# Spell correction



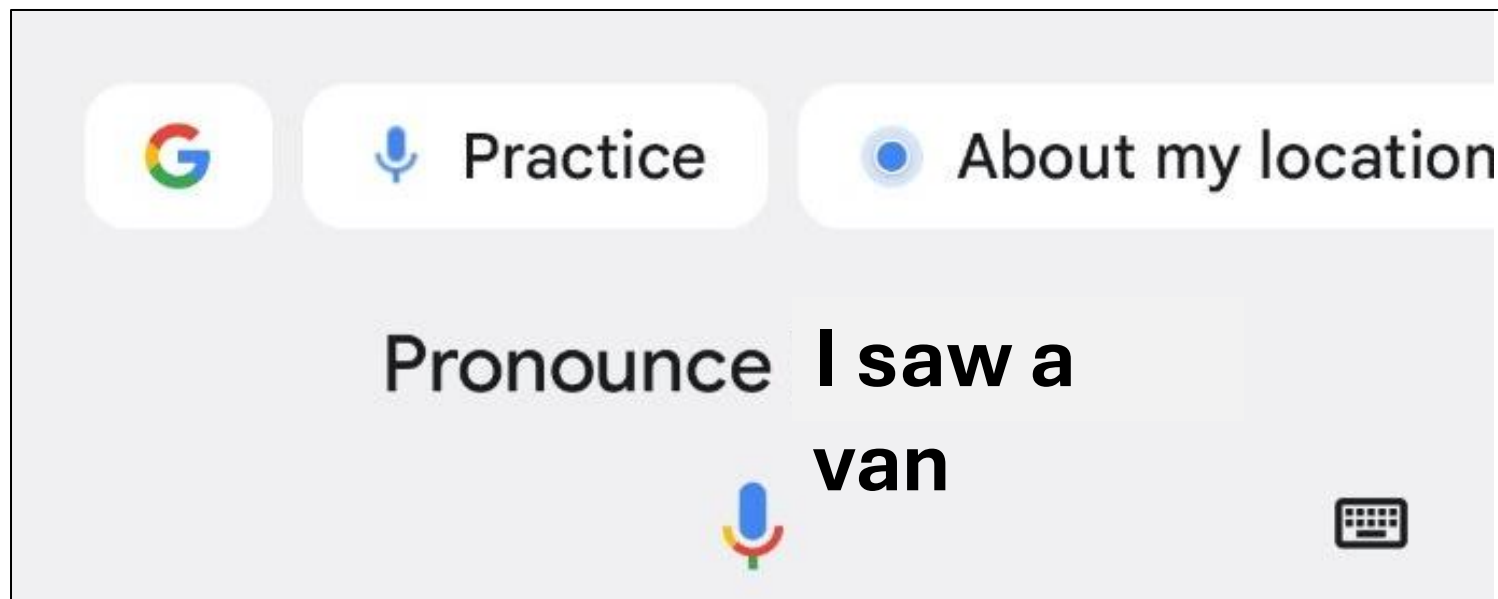
- The office is about fifteen minuets from my house
  - $P(\text{about fifteen minutes from}) > P(\text{about fifteen minuets from})$

# Machine Translation: Important Things To Know



$P(\text{high winds tonite}) > P(\text{large winds tonite})$

# Speech recognition



Here,  $P(I \text{ saw a van}) \gg P(\text{eyes awe of an})$

# Types of Language Models

- Probabilistic language models (PLMs)
- Neural language models (NLMs)

# Probabilistic Language Modeling

- **Goal:** compute the probability of a *sentence* or *sequence of words*:

$$P(W) = P(w_1, w_2, w_3, w_4, w_5, \dots, w_n)$$

- Related task: probability of an upcoming word:

$$P(w_5 \mid w_1, w_2, w_3, w_4)$$

- A model that computes either of these:

$$P(W) \text{ or } P(w_n \mid w_1, w_2, \dots, w_{n-1}) \text{ is called a language}$$

model.

# How to compute $P(W)$

- How to compute this joint probability:
  - $P(\textit{its}, \textit{water}, \textit{is}, \textit{so}, \textit{transparent}, \textit{that})$
- Intuition: let's rely on the Chain Rule of Probability

# Reminder: The Chain Rule

- Recall the definition of conditional probabilities

$$P(B | A) = P(A,B) / P(A)$$

*Rewriting:*  $P(A,B) = P(A)P(B | A)$

- More variables:

$$P(A,B,C,D) = P(A)P(B | A)P(C | A,B)P(D | A,B,C)$$

- The Chain Rule in General

$$P(x_1, x_2, x_3, \dots, x_n) = P(x_1)P(x_2 | x_1)P(x_3 | x_1, x_2) \dots P(x_n | x_1, \dots, x_{n-1})$$

# Chain Rule

*To Compute joint probability of words in a sentence*

*"its water is so transparent"*

$$P(w_1 w_2 \dots w_n) = \prod_i P(w_i | w_1 w_2 \dots w_{i-1})$$

$P(\text{"its water is so transparent"}) =$

$P(\text{its}) \times P(\text{water} | \text{its}) \times P(\text{is} | \text{its water})$

$\times P(\text{so} | \text{its water is}) \times P(\text{transparent} | \text{its water is so})$



# How to estimate these probabilities

Could we just count and divide?

$P(\text{the |its water is so transparent that}) =$   
 $\text{Count(its water is so transparent that the)}$   
 $\text{Count(its water is so transparent that)}$

*Why Not?*

*No!!!! Too many possible sentences!*  
*We will never see enough data to estimate these*



# Markov Assumption :

- Simplifying assumption:

$$P(\text{the l its water is so transparent that}) \approx P(\text{the l that})$$

- Or maybe

$$P(\text{the l its water is so transparent that}) \approx P(\text{the l transparent that})$$

# Markov Assumption

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i \mid w_{i-k} \dots w_{i-1})$$

- In other words, we approximate each component in the product

$$P(w_i \mid w_1 w_2 \dots w_{i-1}) \approx P(w_i \mid w_{i-k} \dots w_{i-1})$$

# Simplest case: Unigram model

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i)$$

Some automatically generated sentences from a unigram model

fifth, an, of, futures, the, an, incorporated, a,  
a, the, inflation, most, dollars, quarter, in, is,  
mass

thrift, did, eighty, said, hard, 'm, july, bullish

that, or, limited, the

# Bigram model

- Condition on the previous word:

$$P(w_i | w_1 w_2 \dots w_{i-1}) \approx P(w_i | w_{i-1})$$

texaco, rose, one, in, this, issue, is, pursuing, growth, in,  
a, boiler, house, said, mr., gurria, mexico, 's, motion,  
control, proposal, without, permission, from, five, hundred,  
fifty, five, yen

outside, new, car, parking, lot, of, the, agreement, reached

this, would, be, a, record, november

# Estimating bigram probabilities

The Maximum Likelihood Estimate

$$P(w_i \mid w_{i-1}) = \frac{\textit{count}(w_{i-1}, w_i)}{\textit{count}(w_{i-1})}$$

$$P(w_i \mid w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

# An example

$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

# Solution?



# Solution

$$P(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

<s> I am Sam </s>

<s> Sam I am </s>

<s> I do not like green eggs and ham </s>

$$P(\text{I} | \text{<s>}) = \frac{2}{3} = .67$$

$$P(\text{Sam} | \text{<s>}) = \frac{1}{3} = .33$$

$$P(\text{am} | \text{I}) = \frac{2}{3} = .67$$

$$P(\text{</s>} | \text{Sam}) = \frac{1}{2} = 0.5$$

$$P(\text{Sam} | \text{am}) = \frac{1}{2} = .5$$

$$P(\text{do} | \text{I}) = \frac{1}{3} = .33$$

# Raw bigram counts

- Total- 9333 sentences

	i	want	to	eat	chinese	food	lunch	spend
i	5	827	0	9	0	0	0	2
want	2	0	608	1	6	6	5	1
to	2	0	4	686	2	0	6	211
eat	0	0	2	0	16	2	42	0
chinese	1	0	0	0	0	82	1	0
food	15	0	15	0	1	4	0	0
lunch	2	0	0	0	0	1	0	0
spend	1	0	1	0	0	0	0	0

Normalize by unigrams:

i	want	to	eat	chinese	food	lunch	spend
2533	927	2417	746	158	1093	341	278

Result:

	i	want	to	eat	chinese	food	lunch	spend
i	0.002	0.33	0	0.0036	0	0	0	0.00079
want	0.0022	0	0.66	0.0011	0.0065	0.0065	0.0054	0.0011
to	0.00083	0	0.0017	0.28	0.00083	0	0.0025	0.087
eat	0	0	0.0027	0	0.021	0.0027	0.056	0
chinese	0.0063	0	0	0	0	0.52	0.0063	0
food	0.014	0	0.014	0	0.00092	0.0037	0	0
lunch	0.0059	0	0	0	0	0.0029	0	0
spend	0.0036	0	0.0036	0	0	0	0	0

# Bigram estimates of sentence probabilities

$$\begin{aligned} P(<s> \text{ I want english food } </s>) = \\ & P(\text{I} | <s>) \\ & \times P(\text{want} | \text{I}) \\ & \times P(\text{english} | \text{want}) \\ & \times P(\text{food} | \text{english}) \\ & \times P(</s> | \text{food}) \\ & = .000031 \end{aligned}$$

# What kinds of knowledge?

- $P(\text{english} | \text{want}) = .0011$
- $P(\text{chinese} | \text{want}) = .0065$
- $P(\text{to} | \text{want}) = .66$
- $P(\text{eat} | \text{to}) = .28$
- $P(\text{food} | \text{to}) = 0$
- $P(\text{want} | \text{spend}) = 0$
- $P(i | \langle s \rangle) = .25$

# Practical Issues


- We do everything in log space
  - Avoid underflow
  - (also adding is faster than multiplying)

$$\log(p_1 \times p_2 \times p_3 \times p_4) = \log p_1 + \log p_2 + \log p_3 + \log p_4$$

# Summary- Language Modeling Task

**Language Modeling** is the task of predicting what word comes next

the students opened their \_\_\_\_\_



- More formally: given a sequence of words , compute the probability distribution of the next word
- A system that does this is called a **Language Model**

# Summary: N-gram Language Models

the students opened their \_\_\_\_\_

- Question: How to learn a Language Model?
- Answer (pre- Deep Learning): learn an n-gram Language Model!
  - Definition: An n-gram is a chunk of n consecutive words.
  - **unigrams**: “the”, “students”, “opened”, “their”
  - **bigrams**: “the students”, “students opened”, “opened their”
  - **trigrams**: “the students opened”, “students opened their”
  - **four-grams**: “the students opened their”
- Idea: *Collect statistics about how frequent different n-grams are and use these to predict next word.*

# N-gram models

- Extend to trigrams, 4-grams, 5-grams
- In general, this is an insufficient model of language because language has **long-distance dependencies**:

*“The computer which I had just put into the machine room on the fifth floor crashed.”*

- But we can often get away with N-gram models



# Smoothing in N-grams

- Like many statistical models, the N-gram probabilistic language model is dependent on the training corpus.
- One practical issue with this is that some word sequences and phrases appear in practice (or in the test set), may not also occur in the training set.
- **SPARSITY AND STORAGE PROBLEMS**
- Important to train robust models that generalize well to handle the **unseen words** and **zero probabilities**



# Add one Smoothing

- Also called **Laplace smoothing**
- Pretend we saw each word one more time than we did
- Just add one to all the counts!

MLE estimate:

$$P_{MLE}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i)}{c(w_{i-1})}$$

Add-1 estimate:

$$P_{Add-1}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i) + 1}{c(w_{i-1}) + V}$$

# Recall the Language modeling task

**Input:** sequence of words,  $x^{(1)}, x^{(2)}, x^{(3)}, x^{(4)}, \dots, x^{(t)}$

**Output:** probability distribution of the next word:  $P(x^{(t+1)} | x^{(t)}, \dots, x^{(1)})$

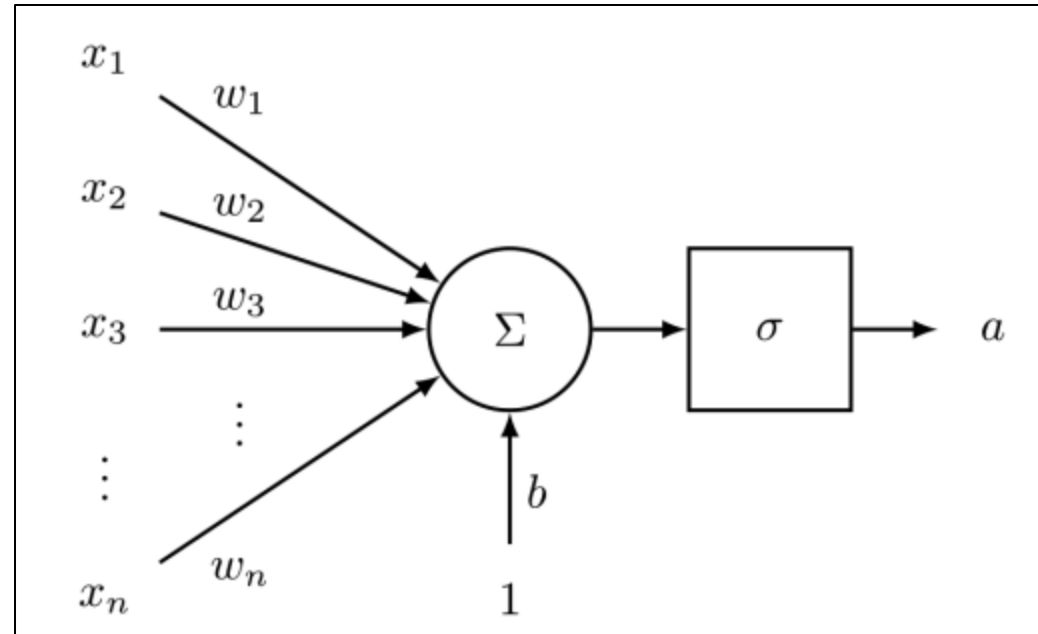
How about a window-based neural model?

# Before building Neural language models


Let's first look into basics of Neural Network.....

# Neurons

A neuron is the fundamental building block of neural networks



# Neural Networks/Feed- Forward Neural Network



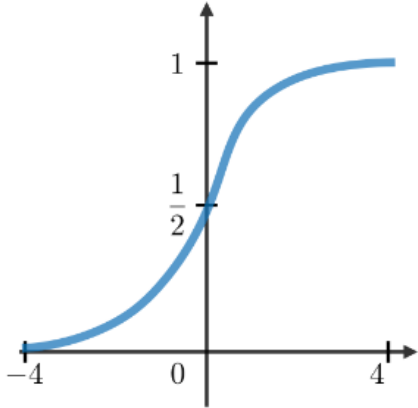
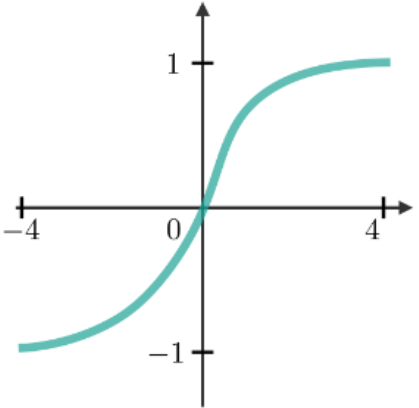
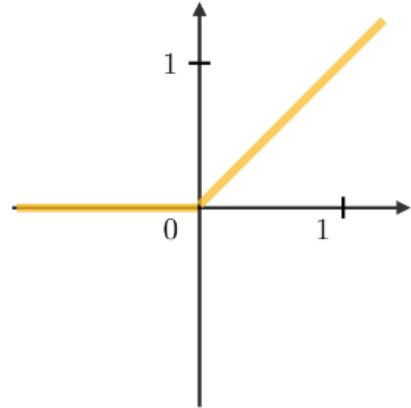
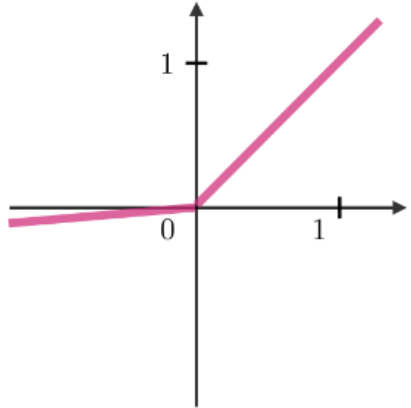
$$z_j^{[i]} = w_j^{[i]T} x + b_j^{[i]}$$

where we denote,  $w$ ,  $b$ ,  $z$  as the weight, bias and output respectively.

# Activation functions

Activation functions are used at the end of a hidden unit to introduce non-linear complexities to the model.

Here are the most common ones:

Sigmoid	Tanh	ReLU	Leaky ReLU
$g(z) = \frac{1}{1 + e^{-z}}$	$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	$g(z) = \max(0, z)$	$g(z) = \max(\epsilon z, z)$ with $\epsilon \ll 1$
			

# Training Neural Networks

*In a neural network, weights are updated as follows:*

Step 1: Take a batch of training data.

Step 2: Perform forward propagation to obtain the corresponding loss.


Step 3: Backpropagate the loss to get the gradients.

Step 4: Use the gradients to update the weights of the network.



# Neural Network Loss function (Cross-entropy loss)

$$-\left[y \log(z) + (1 - y) \log(1 - z)\right]$$



# Learning rate

- **Learning rate:**  $\alpha$  or sometimes  $\eta$ , indicates at which pace the weights get updated.
- This can be fixed or adaptively changed.
- The current most popular method is called Adam, which is a method that adapts the learning rate.

# Backpropagation

- Backpropagation is a method to update the weights in the neural network by taking into account the actual output and the desired output.
- The derivative with respect to weight  $w$  is computed using **chain rule** and is of the following form:

$$\frac{\partial L(z, y)}{\partial w} = \frac{\partial L(z, y)}{\partial a} \times \frac{\partial a}{\partial z} \times \frac{\partial z}{\partial w}$$

- Weight is updated as follows:

$$w \leftarrow w - \alpha \frac{\partial L(z, y)}{\partial w}$$

# Backpropagation Visualization

**Chain Rule!!!!**

[https://docs.google.com/presentation/d/1pmstGvQColwIDP9fJkVLDlG4BNRIfJu8cJTDcokUjrM/edit#slide=id.g27be483e10\\_0\\_0](https://docs.google.com/presentation/d/1pmstGvQColwIDP9fJkVLDlG4BNRIfJu8cJTDcokUjrM/edit#slide=id.g27be483e10_0_0)

# Output of Neural Networks

- *For example: Classification task/Text classification*
- *Neural networks* are capable of producing raw output scores for each of the classes.
- How do we convert output scores into probabilities?

# Interpreting logits: Sigmoid

Sigmoid function:  $\sigma: \mathbb{R} \rightarrow [0, 1]$

$$\sigma(z) = e^z / 1 + e^z = 1 / (1 + e^{-z})$$

- **Class A** (also called the positive class)
- **Not Class A** (complement of Class A or also called the negative class)

# Interpreting logits: SoftMax

Presenting the **softmax** function  $S : \mathbf{R}^C \rightarrow [0, 1]^C$

$$S(\mathbf{z})_i = \frac{e^{\mathbf{z}_i}}{\sum_{j=1}^C e^{\mathbf{z}_j}} = \frac{e^{\mathbf{z}_i}}{e^{\mathbf{z}_1} + \dots + e^{\mathbf{z}_j} + \dots + e^{\mathbf{z}_C}}$$

This function takes a vector of real-values and converts each of them into corresponding probabilities. In a  $C$ -class classification where  $k \in \{1, 2, \dots, C\}$ , it naturally lends the interpretation


$$\text{prob}(y = k | \mathbf{x}) = \frac{e^{\mathbf{z}(\mathbf{x})_k}}{\sum_{j=1}^C e^{\mathbf{z}(\mathbf{x})_j}}$$

# SoftMax Classifier Implementation

[https://docs.google.com/presentation/d/1dFbY-Buku18xhHHbFVNPNuc6vsPeR-LbE8KQXzvuzZg/edit#slide=id.g27be483e10\\_0\\_0](https://docs.google.com/presentation/d/1dFbY-Buku18xhHHbFVNPNuc6vsPeR-LbE8KQXzvuzZg/edit#slide=id.g27be483e10_0_0)




# Binary Classification/Single Label



**Mutually exclusive and exhaustive**, i.e., an input instance can belong to either class, but not both and their **probabilities sum to 1**.


# Multi-class Classification/Single Label



**Mutually exclusive and exhaustive**, i.e., an input instance can belong to only one of these classes, not more and **their probabilities sum to 1**.

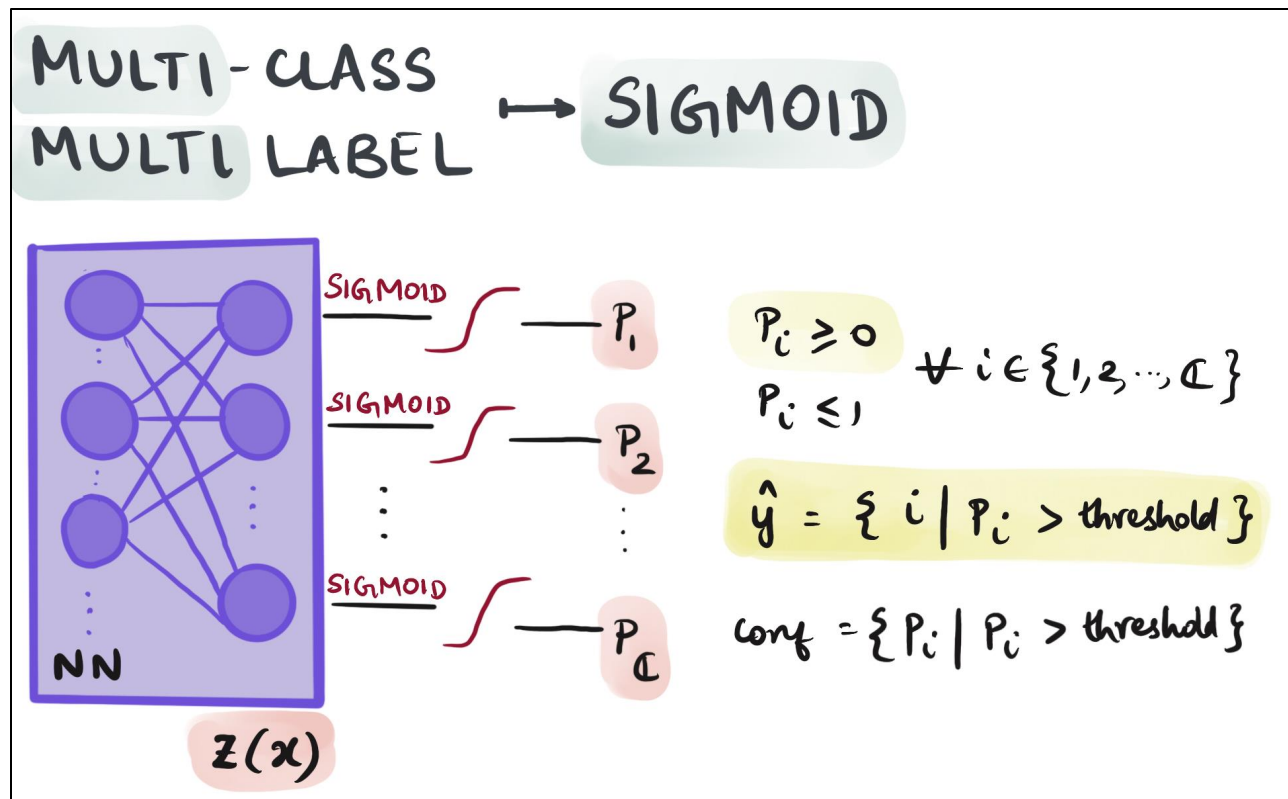
# Multi-class Classification/Multi Label

If input data can belong to more than one class in a multi-class classification problem?



For instance, Genre  
classification of movies (**a  
movie can fall into multiple  
genres**) or classification of  
chest x-rays (**a given chest x-  
ray can have more than one  
disease**).

# Multi-class Classification/Multi Label



Here Classes are NOT mutually exclusive. i.e., train a binary classifier independently for each class. This can be done easily by just applying **sigmoid function** to each of raw scores.

**Note that the output probabilities will NOT sum to 1.**

# Implementation of SoftMax/Sigmoid


```
import torch

def getSoftmaxScores(inputs, dimen):
    ''' Get the softmax scores '''
    print('---Softmax---')
    print('---Dim = ' + str(dimen) + '---')
    softmaxFunc = torch.nn.Softmax(dim = dimen)
    softmaxScores = softmaxFunc(inputs)
    print('Softmax Scores: \n', softmaxScores)
    sums_0 = torch.sum(softmaxScores, dim=0)
    sums_1 = torch.sum(softmaxScores, dim=1)
    print('Sum over dimension 0: \n', sums_0)
    print('Sum over dimension 1: \n', sums_1)

def getSigmoidScores(inputs):
    ''' Get the sigmoid scores: they are element-wise '''
    print('---Sigmoid---')
    sigmoidScores = torch.sigmoid(inputs)
    print('Sigmoid Scores: \n', sigmoidScores)

logits = torch.randn(2, 3)*10 - 5
print('Logits: ', logits)
```

# A fixed-window Neural Language model



# A fixed-window Neural Language Model

Approximately: Y. Bengio, et al. (2000/2003): A Neural Probabilistic Language Model


**Improvements** over  $n$ -gram LM:

- No sparsity problem
- Don't need to store all observed  $n$ -grams

Remaining **problems**:

- Fixed window is **too small**
  - Enlarging window enlarges  $W$
  - Window can never be large enough!
  - $x^{(1)}$  and  $x^{(2)}$  are multiplied by completely different weights in  $W$ .
- No symmetry** in how the inputs are processed.

We need a neural architecture  
that can process *any length* input



# References

- [1] <https://web.stanford.edu/class/archive/cs/cs224n/cs224n.1234/slides/cs224n-2023-lecture05-rnnlm.pdf>
- [2] [https://web.stanford.edu/~jurafsky/slp3/slides/LM\\_4.pdf](https://web.stanford.edu/~jurafsky/slp3/slides/LM_4.pdf)
- [3] <https://stanford-cs324.github.io/winter2022/lectures/introduction/>



# Acknowledgments

- These slides were adapted from the book SPEECH and LANGUAGE PROCESSING: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition
- Practical Natural Language Processing (A Comprehensive Guide to Building Real-World NLP Systems) O'reilly and some modifications from presentations and resources found in the WEB by several scholars.